



IIT Roorkee



Advanced Certification in Cloud Computing & DevOps

VM Migration



- “Live Migration of Virtual Machines”, Christopher Clark, Keir Fraser, Steven Hand, Jacob Gorm Hansen, Eric Jul, Christian Limpach, Ian Pratt, Andrew Warfield
- “Post-Copy Based Live Virtual Machine Migration Using Adaptive Pre-Paging and Dynamic Self-Ballooning”, Michael R. Hines and Kartik Gopalan

Source Contents Credit to:




IIT Roorkee



- **Prof. Rajkumar Buyya** (*University of Melbourne, Australia.*)
- **Prof. Mythili Vutukuru** (IIT Bombay, *India.*)

VM Migration



Review: Virtual Machines

Virtualization provides **interface identical to underlying bare hardware for each VM**

- i.e., all devices, interrupts, memory, page tables etc.

Virtualization Software: VMWare, Hyper-V, Oracle virtual box, Xen etc.

- It allows clean **separation between hardware and software**.
- **Process level migration problems can be avoided** by migrating a virtual machine.
- **Virtualization provides facility to migrate** virtual machine from one host (source) to another physical host (destination).
- Virtual Machine Migration is a **useful tool for administrator** of data center and clusters.

Review: Virtual Machines in Cloud

☐ Benefits of Virtual Machines

- ☐ Virtualization helps in making **efficient use of hardware resources**
- ☐ Facilitates a greater **degree of abstraction**
- ☐ **Easily move** from one piece of hardware to another
- ☐ **Replicate** them at will
- ☐ Create more **scalable and flexible infrastructure**
- ☐ **High Availability and Fault tolerance (How?)**

☐ Cloud computing has taken that **degree of efficiency and agility realized** from virtualization


- ☐ Pooled resources
- ☐ Geographic diversity
- ☐ Universal connectivity

Motivation

- Consider a **data center** consisting of “**one**” **physical machines (PM)** hosting “**two**” **VMs** implementing one customer application each
- Resources**(CPU, Network, Memory, I/O) are **allocated to each VM** to **handle the workload** and operate at certain performance level (**SLA**)
- Each **VM sees workload fluctuation (WHY?)** from time to time =>

resource requirement changes

of user visit
increases



PM
Capacity

Network Bandwidth
Memory
CPU

= 10 Gbps
= 16 GB
= 8 cores

Resource
Allocation

VM1

N = 5Gbps
M = 8GB
C = 4 cores

VM2

N = 5Gbps
M = 8GB
C = 4 cores

- An **increase in workload** can be handled by **allocating more resources** to it, if **idle** resources are available

❑ Main Issues:

What if PM does not have (enough or no) idle resources to satisfy VM's requirement?

- **Performance** of the application degrades
- **SLA violation** occurs

❑ Key Ideas

- Replication VMs
- Migrating VMs

Virtual Machine Migration

- Why do we need migration?
- When do we need to migrate?
- How migration is done?
- Issues in long distance migration (across data centers)

• Why do we need migration?

- To Maintain/Upgrade
- To Optimize the Performance
- To Protect SLA
- To Improve QoE

When we need to migrate? [NSDI' 07]

- **Hotspots can cause SLA violations**

- Burden on some Virtual or Physical Machines are called hotspots

- **Hotspot Detection**

- ***Black-box Monitoring***

- ☐ CPU

- ☐ Network

- ☐ Memory

- ***Gray-box Monitoring***

- ☐ Gather OS level statistics and application logs

- A hotspot is **flagged only if thresholds or SLAs are exceeded** for a sustained time

When we need to migrate? [FGCS' 12]

- **SLA violation detection**

- **Mapping** low-level resource metrics to high-level SLAs
 - *CPU speed maps to Response Time*
 - *Occupied memory size maps to number of concurrent clients*
- **Predictive Strategy** for detection of possible SLA violations
- **Detection interval**
 - **Short measurement intervals** may **degrade** performance
 - **Long measurement intervals** may cause **ignorance of heavy SLA** violations.

Migration Techniques

- The **different virtual machine migration** techniques are as follows:
 - Fault Tolerant Migration Techniques
 - Load Balancing Migration Techniques
 - Energy Efficient Migration Techniques

Fault Tolerant Migration Techniques

- Fault tolerance **allows the VMs to continue its job even any part of system fails.**
- This technique **migrates the VM from one physical server to another physical server based upon the prediction of the failure occurred.**
- Fault tolerant migration technique is **to improve the availability of physical server and avoids performance degradation of applications.**

- The Load balancing migration technique **aims to distribute load across the physical servers to improve the scalability of physical servers** in cloud environment.
- **The Load balancing improves the**
 - resource consumption
 - implementation of failover
 - enhancing scalability

- **The power consumption of Data center is mainly based on the utilization of the servers and their cooling systems.**
- **The servers typically need up to 70 percentage of their maximum power consumption even at their low utilization level.**
- **Therefore, there is a need for migration techniques that conserves the energy of servers by optimum resource utilization.**

VM Migration Methods

- Virtual Machine Migration methods are divided into **two types**:
 - **Cold (non-live) migration**
 - **Hot (live) migration**
- The **status of the VM loses** and **user can notice the service interruption** in cold migration. First, **VM is suspended**, then its **state is transferred**, and lastly, **VM is resumed at destination host**.
- In **live migration process**, the **state of a running VM is transferred/migrated from Host A to Host B**.
- **VM keeps running** while migrating and **does not lose its status**. **User doesn't feel any interruption** in service in hot (live) migration.

Live (Hot) VM Migration

Recall: Live Migration

- VM live migration can be a extremely **powerful tool for cluster/system administrators.**

Hardware / Software maintenance / upgrades

Load balancing / resource management

Distributed power management



- **Move VM instances across distinct physical hosts with **little or no downtime** for running services.**
 - **Services are unaware** of the migration.
 - Maintain **network connections of the guest OS**.
 - **VM is treaded as a black box**.
 - **VM (OS level) Migration is easier than migrating processes**




Why OS-level migration, instead of process-level?

- *Avoid 'residual dependencies'*
 - Original host can be power-off / sleep once migration completed.
- *Can transfer in-memory state in a consistent and efficient fashion*
 - E.g., No reconnection for media streaming application
- *Allow a separation of concerns b/w the users and operator of a cluster*
 - Users can fully control of the software and services within their VM.
 - Operators don't care about what's occurring within the VM.

What is migrated?

- **CPU context of VM, contents of main memory**
 - Narrow interface, easier than process migration
- **Disk: assume NAS (network attached storage) that is accessible from both hosts, or local disk is mirrored**
 - We do not consider migrating disk data
- **Network: assume both hosts on same LAN**
 - Migrate IP address, advertise new MAC address to IP mapping via ARP reply
 - Migrate MAC address, let switches learn new MAC location
 - Network packets redirected to new location (with transient losses)
- **I/O devices are provisioned at target**
 - Virtual I/O devices easier to migrate.

Design-local resources



Steps to Migrate a VM

- Broad steps in any migration technique: **Suppose** we are migrating a VM from host A to host B
 1. **Setup target host B, reserve resources for the VM**
 2. **Push phase:** push some memory of VM from A to B
 3. **Stop-and-copy:** stop the VM at A, copy CPU context, and some memory
 4. **Pull phase:** Start VM at host B, pull any further memory required from A
 5. **Clean up state from host A, migration complete**
- **Total migration time:** time for steps 2,3,4
- **Service downtime:** time for step 3
- **Other metrics:** impact on application performance, network bandwidth consumed, total pages transferred

Flavors of migration techniques




IIT Roorkee



- **Pure stop-and-copy**: VM stopped, **all state transferred to target**, VM restarts
 - Too much downtime to be classified as “live” migration
- **Pre-copy**: **most state is transferred in the push phase**, followed by a brief stop-and-copy phase
- **Post-copy**: VM stopped, **bare minimum state required to run the VM is transferred to the target host. Remaining state is pulled on demand** while the VM is running at the new location.
- **Hybrid**: **a mix of pre-copy and post-copy**. Some pushing of state followed by stop-and-copy, followed by pulling of state on demand.

Design-challenges

- ☐ Minimize service downtime
- ☐ Minimize migration duration
- ☐ Avoid disrupting running service



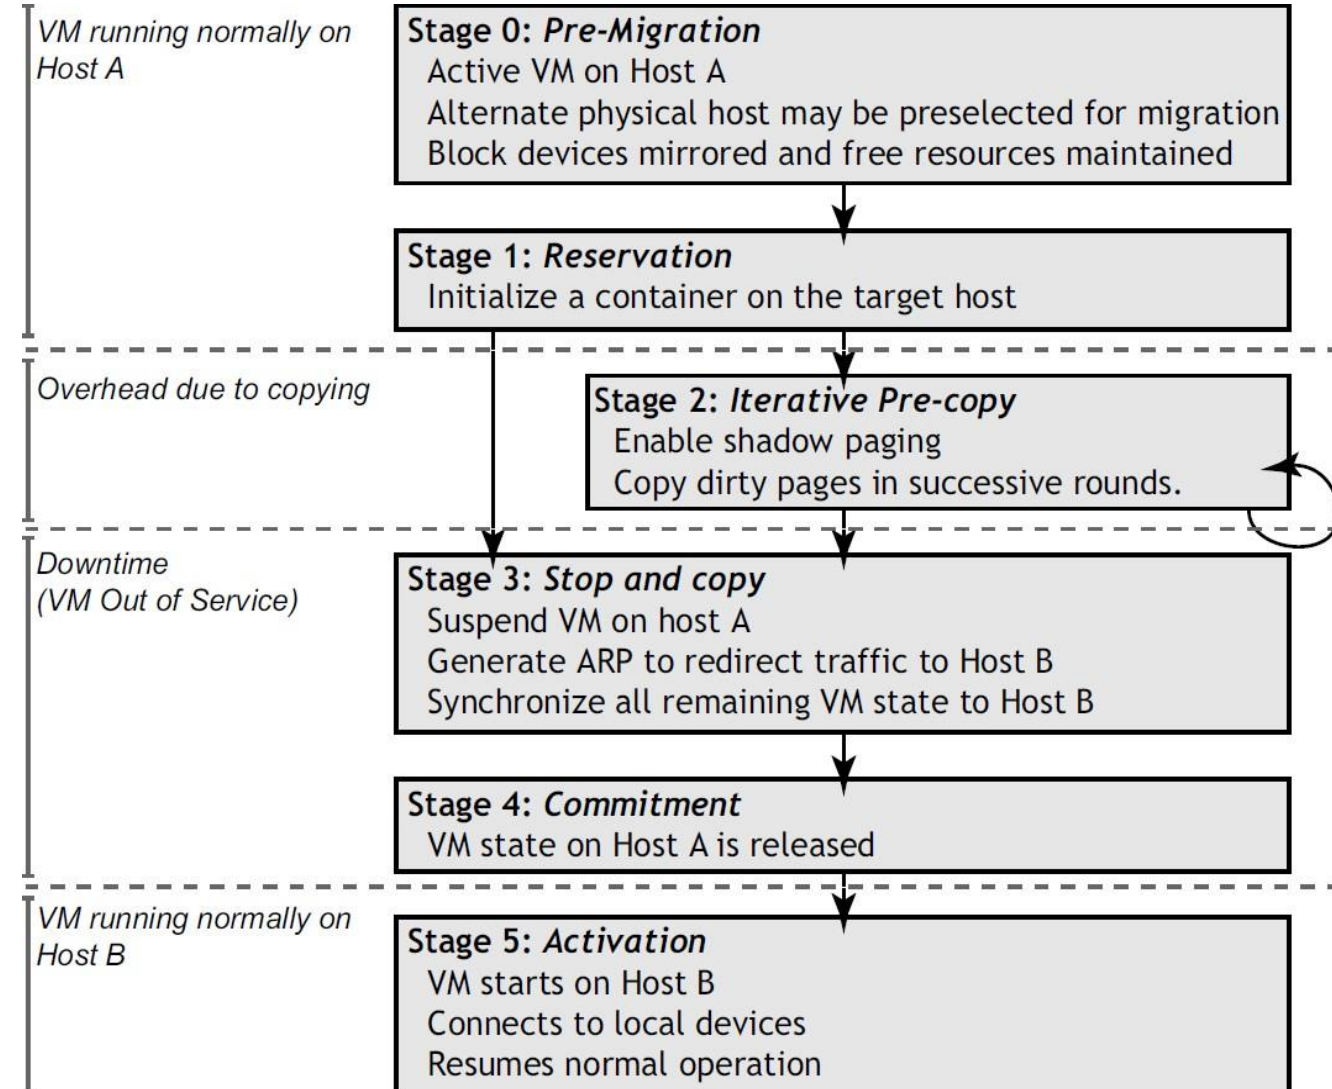
Design-memory migration

Phase	service downtime	migration duration
push	-	-
stop-and-copy	longest	shortest
pull (demand)	shortest	longest

- Careful to avoid service degradation


Pre-copy based livemigration

- Iterative pre-copy + stop-and-copy for remaining memory
- First push round copies all pages
- Every round copies pages dirtied in previous round
 - A page maybe copied multiple times
- **Writable Working Set (WWS):** pages commonly written to
 - WWS will be copied multiple times
 - Finally transferred in stop-and-copy
- **How many rounds?**
- Stop when rate of dirtying $>$ rate of transfer




Impact of iterative pre-copy

Effect of Bandwidth and Pre-Copy Iterations on Migration Downtime
(Based on a page trace of OLTP Database Benchmark)



- If stop-and-copy, 512MB VM, 128 Mbps network, downtime = 32 sec
- With one pre-copy round, downtime goes to 2-3 sec
 - ~1 second for 2 or more rounds



Implementation-dynamic rate limiting

More network bandwidth, less service downtime !

Less network bandwidth, less impact on running service !

Dynamically adapt the bandwidth limit during each round

- Set a minimum and a maximum bandwidth limit, begin with the minimum limit


- $bandwidth_{next} = dirty\ rate_{current} + constant\ increment$

- $dirty\ rate_{current} = \frac{dirty\ pages}{duration}$

When terminate push, and switch to stop-and-copy ?

- $dirty\ rate_{current} > bandwidth_{max}$

- $dirty\ pages < threshold$




Some optimizations

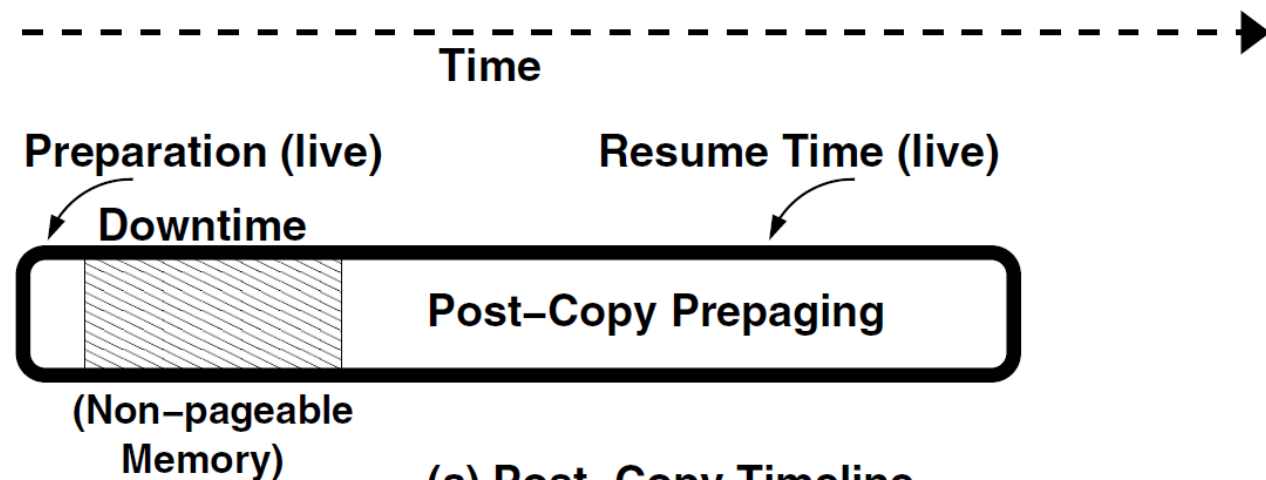
- Avoid transferring page multiple times
 - Before transmitting page, peek into the current round's dirty bitmap
 - Skip transmission if page is already dirtied in ongoing round
- Move non-interactive processes generating dirty pages to wait queue
 - Execution paused until migration completes
- Free up page cache and other unnecessary pages
 - Reduce memory footprint
 - Much like ballooning

Pre-copy vs Post-copy livemigration

- Avoid multiple transfers of same page as happens in pre-copy
- Prepare target, stop VM, copy CPU context and minimum memory to target
- Start VM at target, pull memory from source via demand paging
 - Memory access at target causes page fault, page fetched from source



(a) Pre-Copy Timeline



(a) Post-Copy Timeline

- **Active pushing**: source proactively pushes important pages, in addition to pulling pages via page faults
- **Pre-paging**: a “bubble” of pages around faulted page and proactively pushed, in anticipation of future accesses
- **Dynamic self-ballooning**: VM periodically frees up unnecessary memory and gives it back to hypervisor
 - Reduces memory footprint, speeds up page transfer
 - Performed carefully without hurting application performance
 - Can be used to optimize pre-copy migration as well
- **Hybrid**: one pre-copy round, followed by post copy

Post copy vs pre copy performance

- Longer downtime as compared to pre-copy, but lower total migration time, fewer page transfers, lesser disruption to application




Figure 9. Impact of post-copy on NetPerf bandwidth.


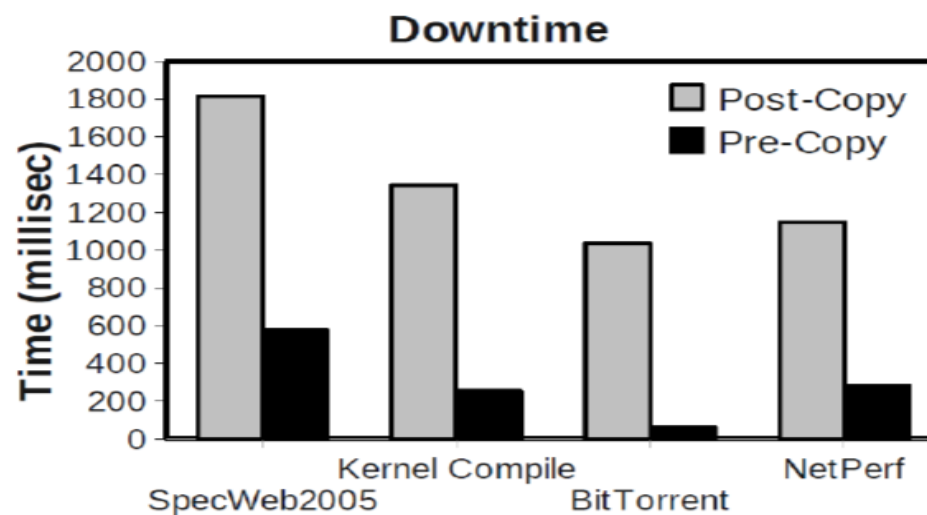
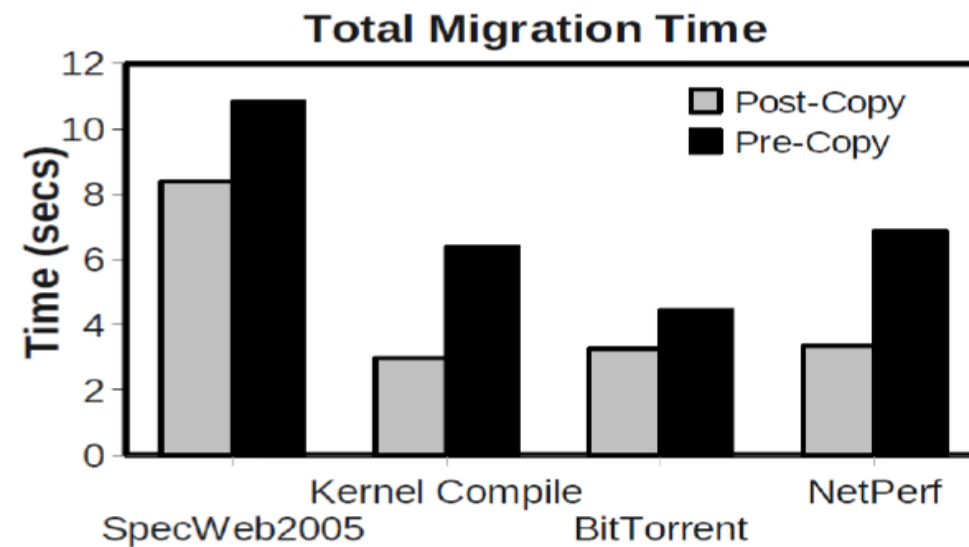



Figure 10. Impact of pre-copy on NetPerf bandwidth.

Post copy vs pre copy performance



What about failures?

- What if **target machine fails during migration**?
- **Pre-copy can simply abort the migration**, restart with another target
 - With pre-copy, **latest state is on source only**, so can recover
- With post copy, **source has stale memory**, target has updated memory
 - If target crashes during post copy, cannot recover application data (unless some replication is performed)

Summary

- VM live migration techniques
 - Iterative pre-copy vs post-copy via demand paging
 - Implementation details on Xen
 - Performance comparison
- Dynamically adapting network-bandwidth
 - Balance service downtime and service performance degradation
- Which is better?
 - Pre-copy suited for interactive application
 - Post copy is better for memory-intensive applications with large WWS
 - Hybrid techniques are also used



IIT Roorkee



India: +91-7022374614

US: 1-800-216-8930 (TOLL FREE)



support@intellipaate.com



24/7 Chat with Our Course Advisor